

# Hive Programming



# Topics Covered

- ◆ About Hive
- ◆ Comparing Hive to SQL
- ◆ Hive Architecture
- ◆ Submitting Hive Queries
- ◆ Defining Tables
- ◆ Loading Data into Hive
- ◆ Performing Queries
- ◆ Hive Partitions, Buckets, and Skewed
- ◆ Sorting Data
- ◆ Hive Join Strategies



■ About Hive




# Objectives



# About Hive

**Store and Query all  
Data in Hive**



**Use Existing SQL Tools  
and Existing SQL Processes**



## About Hive – cont.

- ◆ It is a data warehouse system for Hadoop
- ◆ It maintains metadata information about your big data stored on HDFS
- ◆ It treats your big data as tables
- ◆ It performs SQL-like operations on the data using a scripting language called **HiveQL**



# Objectives



- ◆ About Hive
- ◆ Comparing Hive to SQL





# Hive's Alignment with SQL

## SQL Datatypes

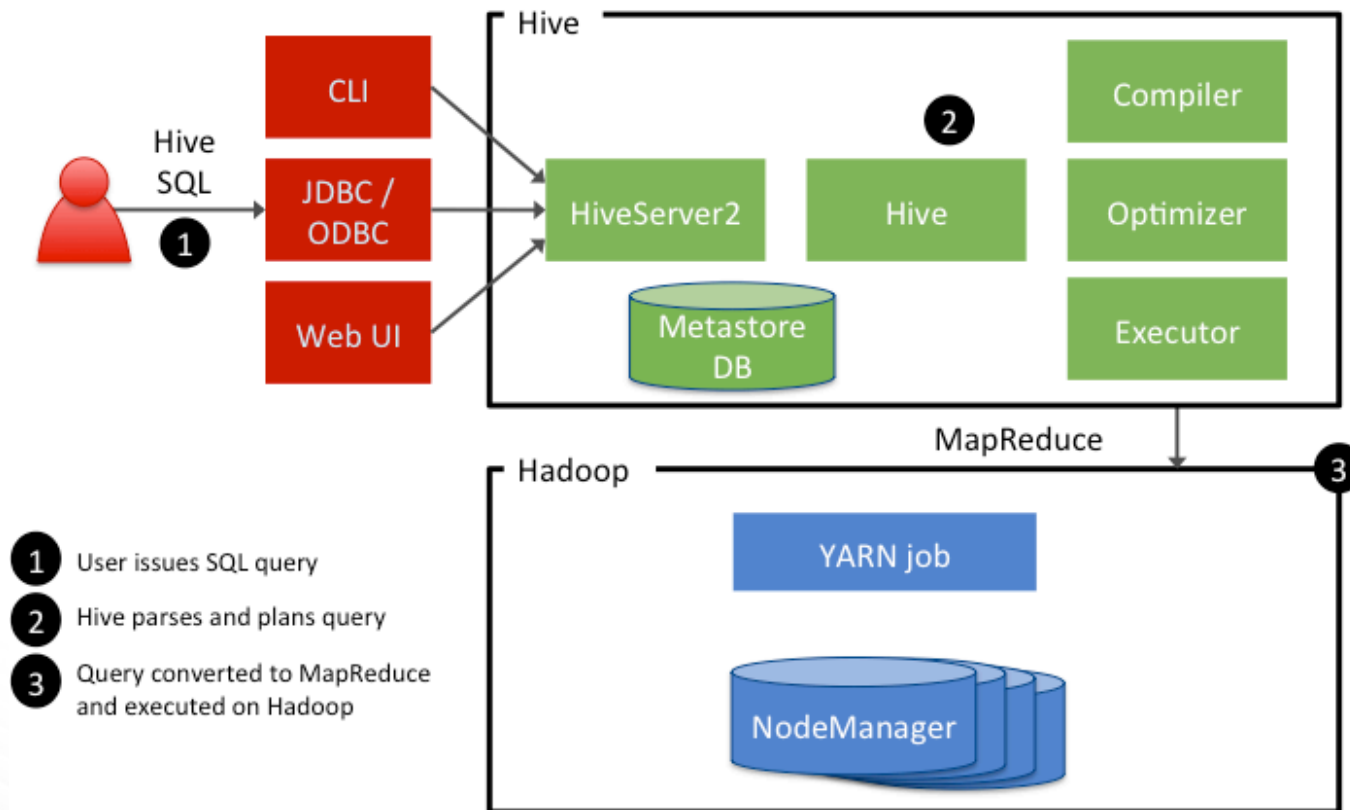
INT
TINYINT/SMALLINT/BIGINT
BOOLEAN
FLOAT
DOUBLE
STRING
BINARY
TIMESTAMP
ARRAY, MAP, STRUCT, UNION
DECIMAL
CHAR
VARCHAR
DATE

## SQL Semantics

SELECT, LOAD, INSERT from query
Expressions in WHERE and HAVING
GROUP BY, ORDER BY, SORT BY
CLUSTER BY, DISTRIBUTE BY
Sub-queries in FROM clause
GROUP BY, ORDER BY
ROLLUP and CUBE
UNION
LEFT, RIGHT and FULL INNER/OUTER JOIN
CROSS JOIN, LEFT SEMI JOIN
Windowing functions (OVER, RANK, etc.)
Sub-queries for IN/NOT IN, HAVING
EXISTS / NOT EXISTS



# Hive Architecture





# Objectives



- ◆ About Hive
- ◆ Comparing Hive to SQL
- ◆ Hive Queries



# Submitting Hive Queries

## ◆ Hive CLI

- Traditional Hive “thick” client

- ```
$ hive
```

```
hive>
```

## ◆ Beeline

- A new command-line client that connects to a HiveServer2 instance

- ```
$ beeline -u url -n username -p password
```

```
beeline>
```



## Defining a Hive-Managed Table

```
CREATE TABLE customer (  
    customerID INT,  
    firstName STRING,  
    lastName STRING,  
    birthday TIMESTAMP  
) ROW FORMAT DELIMITED  
  FIELDS TERMINATED BY ',';
```



## Defining an External Table

```
CREATE EXTERNAL TABLE salaries (  
    gender string,  
    age int,  
    salary double,  
    zip int  
    ) ROW FORMAT DELIMITED  
    FIELDS TERMINATED BY ',';
```



## Defining a Table LOCATION

```
CREATE EXTERNAL TABLE SALARIES (  
    gender string,  
    age int,  
    salary double,  
    zip int  
) ROW FORMAT DELIMITED  
    FIELDS TERMINATED BY ','  
    LOCATION '/user/train/salaries/';
```



## Loading Data into Hive

```
LOAD DATA LOCAL INPATH '/tmp/customers.csv'  
OVERWRITE INTO TABLE customers;
```

```
LOAD DATA INPATH '/user/train/customers.csv'  
OVERWRITE INTO TABLE customers;
```

```
INSERT INTO TABLE birthdays  
  SELECT firstName, lastName, birthday  
  FROM customers  
  WHERE birthday IS NOT NULL;
```



## Performing Queries

```
SELECT * FROM customers;
```

```
FROM customers  
  SELECT firstName, lastName, address, zip  
  WHERE orderID > 0  
  ORDER BY zip;
```

```
SELECT customers.*, orders.*  
  FROM customers  
  JOIN orders ON  
    (customers.customerID = orders.customerID);
```





# Objectives



- ◆ About Hive
- ◆ Comparing Hive to SQL
- ◆ Hive Queries
- ◆ Hive Partitions, Buckets, and Skewed



# Hive Partitions

- Use the **partitioned by** clause to define a partition when creating a table:


```
create table employees (id int, name string, salary double)
partitioned by (dept string);
```

- Subfolders are created based on the partition values:

```
/apps/hive/warehouse/employees
  /dept=hr/
  /dept=support/
  /dept=engineering/
  /dept=training/
```



# Hive Buckets



## Skewed Tables

```
CREATE TABLE Customers (  
    id int,  
    username string,  
    zip int  
)  
SKEWED BY (zip) ON (57701, 57702)  
STORED as DIRECTORIES;
```



# Objectives



- ◆ About Hive
- ◆ Comparing Hive to SQL
- ◆ Hive Queries
- ◆ Hive Partitions, Buckets, and Skewed
- ◆ Sorting Data



## Sorting Data

Hive has two sorting clauses:

- ◆ **order by:** a complete ordering of the data
- ◆ **sort by:** data output is sorted per reducer



## Using Distribute By

```
insert overwrite table mytable  
  select gender,age,salary  
  from salaries  
distribute by age;
```

```
insert overwrite table mytable  
  select gender,age,salary  
  from salaries  
distribute by age  
sort by age;
```





## Storing Results to a File

```
INSERT OVERWRITE DIRECTORY
```

```
' /user/train/ca_or_sd/'
```

```
from names
```

```
  select name, state
```

```
  where state = 'CA'
```

```
  or state = 'SD';
```

```
INSERT OVERWRITE LOCAL DIRECTORY
```

```
' /tmp/myresults/'
```

```
SELECT * FROM bucketnames
```

```
ORDER BY age;
```



## Specifying MapReduce Properties

```
SET mapreduce.job.reduces = 12
```

```
hive -f myscript.hive  
  -hiveconf mapreduce.job.reduces=12
```

```
SELECT * FROM names  
      WHERE age = ${age}  
hive -f myscript.hive -hivevar age=33
```



# Objectives



- ◆ About Hive
- ◆ Comparing Hive to SQL
- ◆ Hive Queries
- ◆ Hive Partitions, Buckets, and Skewed
- ◆ Sorting Data
- ◆ Hive Join Strategies



## Hive Join Strategies


Type	Approach	Pros	Cons
Shuffle Join	Join keys are shuffled using MapReduce, and joins are performed on the reduce side.	Works regardless of data size or layout.	Most resource-intensive and slowest join type.
Map (Broadcast) Join	Small tables are loaded into memory in all nodes, mapper scans through the large table, and joins.	Very fast, single scan through largest table.	All but one table must be small enough to fit in RAM.
Sort-Merge-Bucket Join	Mappers take advantage of co-location of keys to do efficient joins.	Very fast for tables of any size.	Data must be sorted and bucketed ahead of time.



## Shuffle Joins

customer				orders		
first	last	id		cid	price	quantity
Nick	Toner	11911		4150	10.50	3
Jessie	Simonds	11912		11914	12.25	27
Kasi	Lamers	11913		3491	5.99	5
Rodger	Clayton	11914		2934	39.99	22
Verona	Hollen	11915		11914	40.50	10

SELECT \* FROM customer JOIN orders ON customer.id = orders.cid;



## Map (Broadcast) Joins

customer			orders		
first	last	id	cid	price	quantity
Nick	Toner	11911	4150	10.50	3
Jessie	Simonds	11912	11914	12.25	27
Kasi	Lamers	11913	3491	5.99	5
Rodger	Clayton	11914	2934	39.99	22
Verona	Hollen	11915	11914	40.50	10

`SELECT * FROM customer JOIN orders ON customer.id = orders.cid;`

{ id: 11914, { first: Rodger, last: Clayton } }  
{ cid: 11914, { price: 12.25, quantity: 27 },  
cid: 11914, { price: 12.25, quantity: 27 } }



Records are joined during  
the map phase.



## Sort-Merge-Bucket Joins

customer			orders		
first	last	id	cid	price	quantity
Nick	Toner	11911	4150	10.50	3
Jessie	Simonds	11912	11914	12.25	27
Kasi	Lamers	11913	11914	40.50	10
Rodger	Clayton	11914	12337	39.99	22
Verona	Hollen	11915	15912	40.50	10

`SELECT * FROM customer join orders ON customer.id = orders.cid;`

**Distribute and sort by the most common join key.**

```
CREATE TABLE orders (cid int, price float, quantity int)
CLUSTERED BY(cid) SORTED BY(cid) INTO 32 BUCKETS;
```

```
CREATE TABLE customer (id int, first string, last string)
CLUSTERED BY(id) SORTED BY(cid) INTO 32 BUCKETS;
```





## Invoking a Hive UDF

```
ADD JAR /myapp/lib/myhiveudfs.jar;  
CREATE TEMPORARY FUNCTION  
ComputeShipping  
  AS 'hiveudfs.ComputeShipping';  
  
FROM orders SELECT  
  address,  
  description,  
  ComputeShipping(zip, weight);
```



# Advanced Hive Programming



## Topics Covered

- ◆ Performing a Multi-Table/File Insert
- ◆ Understanding Views
- ◆ Defining Views
- ◆ Using Views
- ◆ The OVER Clause
- ◆ Using Windows
- ◆ Hive Analytics Functions
- ◆ *Lab: Advanced Hive Programming*
- ◆ Hive File Formats
- ◆ Hive SerDe



- Performing a Multi-le/File Insert

## Objectives



## Performing a Multi-Table/File Insert

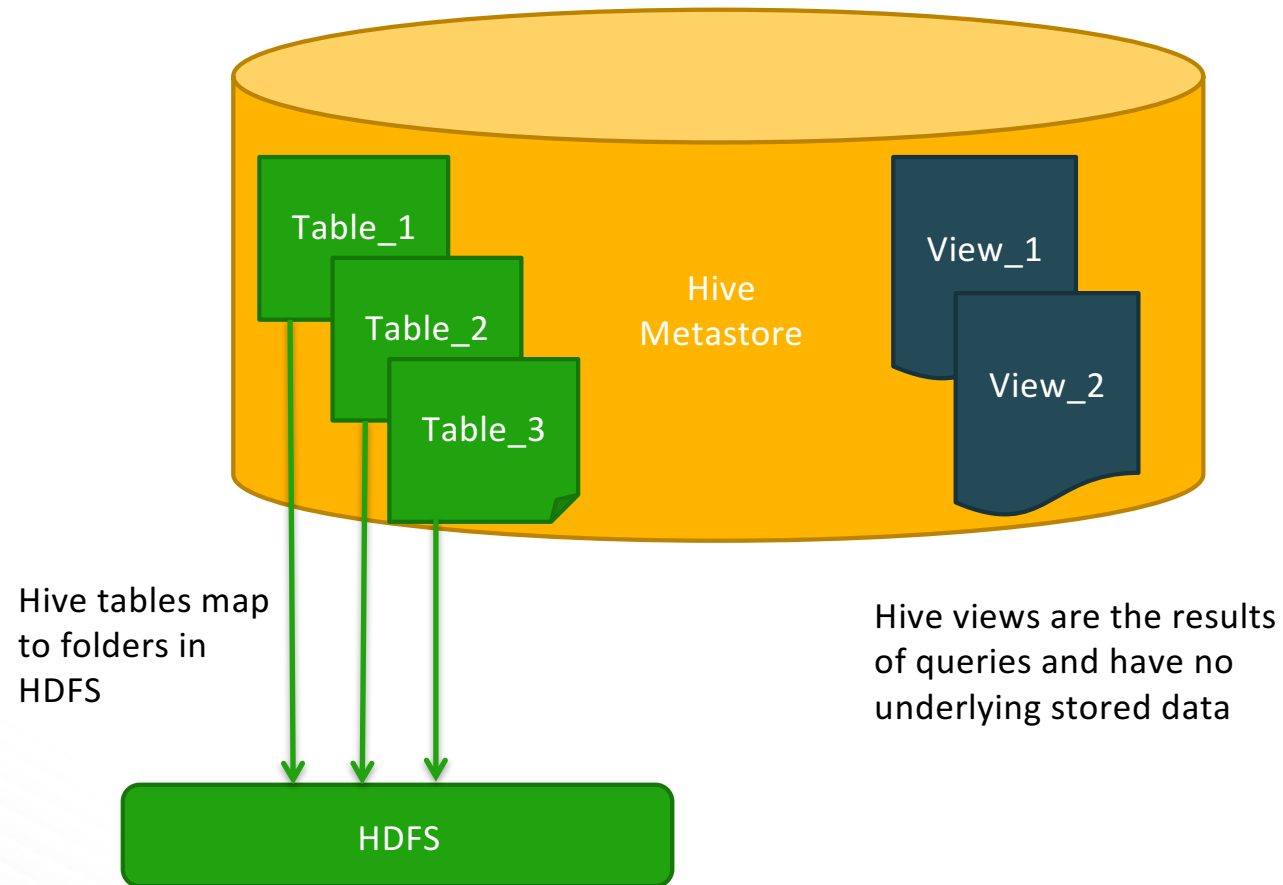
```
insert overwrite directory '2014_visitors' select * from wh_visits  
where visit_year='2014'  
insert overwrite directory 'ca_congress' select * from congress  
where state='CA' ;
```

No semicolon

```
from visitors  
INSERT OVERWRITE TABLE gender_sum  
  SELECT visitors.gender, count_distinct(visitors.userid)  
  GROUP BY visitors.gender  
INSERT OVERWRITE DIRECTORY '/user/tmp/age_sum'  
  SELECT visitors.age, count_distinct(visitors.userid)  
  GROUP BY visitors.age;
```



# Understanding Views





# Objectives

- ◆ Performing a Multi-Table/File Insert
- ◆ Defining and Using Views





## Defining Views

```
CREATE VIEW 2010_visitors AS  
  SELECT fname, lname,  
          time_of_arrival, info_comment  
  FROM wh_visits  
  WHERE  
    cast(substring(time_of_arrival,6,4)  
  AS int) >= 2010  
  AND  
    cast(substring(time_of_arrival,6,4)  
  AS int) < 2011;
```



## Using Views

You use a view just like a table:

```
from 2010_visitors
select *
where info_comment like "%CONGRESS%"
order by lname;
```



## The TRANSFORM Clause

- ◆ You can write your own custom mappers and/or reducers in Hive
- ◆ Use the **TRANSFORM** clause to specify your custom script:

```
add file splitwords.py;
add file countwords.py;

FROM (
  FROM mytable
  SELECT TRANSFORM(keywords) USING 'python splitwords.py'
  AS word, count
  CLUSTER BY word
) wc
INSERT OVERWRITE TABLE word_count
  SELECT TRANSFORM(wc.word, wc.count) USING 'python countwords.py'
  AS word, count;
```



# Objectives



- ◆ Performing a Multi-Table/File Insert
- ◆ Defining and Using Views
- ◆ Using Clauses and Windows



## The OVER Clause

orders				result set	
cid	price	quantity		cid	max(price)
4150	10.50	3	→	2934	39.99
11914	12.25	27		4150	10.50
4150	5.99	5		11914	40.50
2934	39.99	22			
11914	40.50	10			

SELECT cid, max(price) FROM orders GROUP BY cid;

orders				result set	
cid	price	quantity		cid	max(price)
4150	10.50	3	→	2934	39.99
11914	12.25	27		4150	10.50
4150	5.99	5		4150	10.50
2934	39.99	22		11914	40.50
11914	40.50	10		11914	40.50

SELECT cid, max(price) OVER (PARTITION BY cid) FROM orders;



## Using Windows

orders				result set	
cid	price	quantity		cid	sum(price)
4150	10.50	3	→	4150	5.99
11914	12.25	27		4150	16.49
4150	5.99	5		4150	36.49
4150	39.99	22		4150	70.49
11914	40.50	10		11914	12.25
4150	20.00	2		11914	52.75

`SELECT cid, sum(price) OVER (PARTITION BY cid ORDER BY price  
ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) FROM orders;`



## Using Windows – cont.

```
SELECT cid, sum(price) OVER  
(PARTITION BY cid ORDER BY price ROWS  
BETWEEN 2 PRECEDING AND 3 FOLLOWING)  
FROM orders;
```

```
SELECT cid, sum(price) OVER  
(PARTITION BY cid ORDER BY price ROWS  
BETWEEN UNBOUNDED PRECEDING AND  
CURRENT ROW) FROM orders;
```



## Hive Analytics Function

orders				result set		
cid	price	quantity		cid	quantity	rank
4150	10.50	3	→	4150	2	1
11914	12.25	27		4150	3	2
4150	5.99	5		4150	5	3
4150	39.99	22		4150	22	4
11914	40.50	10		11914	10	1
4150	20.00	2		11914	27	2

SELECT cid, quantity, **rank()** OVER (PARTITION BY cid  
ORDER BY quantity) FROM orders;





# Objectives



- ◆ Performing a Multi-Table/File Insert
- ◆ Defining and Using Views
- ◆ Using Clauses and Windows
- ◆ Hive Programming



## Hive File Formats

- ◆ Text file
- ◆ SequenceFile
- ◆ RCFile
- ◆ ORC File

```
CREATE TABLE names  
  (fname string, lname string)  
STORED AS RCFile;
```



# Hive SerDe

- ◆ SerDe = serializer/deserializer
- ◆ Determines how records are read from a table and written to HDFS

```
CREATE TABLE emails (  
  from_field string,  
  sender string,  
  email_body string)  
ROW FORMAT SERDE  
  'org.apache.hadoop.hive.serde2.avro.AvroSerDe'  
STORED AS INPUTFORMAT  
  'org.apache.hadoop.hive ql.io.avro.AvroContainerInputFormat'  
  OUTPUTFORMAT  
  'org.apache.hadoop.hive ql.io.avro.AvroContainerOutputFormat'  
TBLPROPERTIES (  
  'avro.schema.url'='hdfs//nn:8020/emailschema.avsc');
```



## Hive ORC Files

The ***Optimized Row Columnar*** (ORC) file format provides a highly efficient way to store Hive data

```
CREATE TABLE tablename (  
  ...  
) STORED AS ORC;  
  
ALTER TABLE tablename SET FILEFORMAT  
ORC;  
  
SET hive.default.fileformat=Orc
```



## Computing Table and Column Statistics

```
ANALYZE TABLE tablename COMPUTE  
STATISTICS;
```

```
ANALYZE TABLE tablename COMPUTE  
STATISTICS FOR COLUMNS column_name_1,  
column_name_2, ...
```

```
DESCRIBE FORMATTED tablename
```

```
DESCRIBE EXTENDED tablename
```



# Objectives



- ◆ Performing a Multi-Table/File Insert
- ◆ Defining and Using Views
- ◆ Using Clauses and Windows
- ◆ Hive Programming
- ◆ Hive Optimization and Queries



## Hive Cost-Based Optimization (CBO)

- ◆ **Cost-Based Optimization** (CBO) engine uses statistics within Hive tables to produce optimal query plans
- ◆ Two types of stats used for optimization:
  - Table stats
  - Column stats
- ◆ Uses an open-source framework called **Calcite**
- ◆ To use CBO, you need to:
  - Analyze the table and relevant columns
  - Set the appropriate properties



## Optimizing Queries with Statistics

```
analyze table tweets compute statistics;
```

```
analyze table tweets compute statistics for  
columns sender, topic;
```

```
set hive.compute.query.using.stats=true;  
set hive.cbo.enable=true;  
set hive.stats.fetch.column.stats=true;
```





# Vectorization

Before



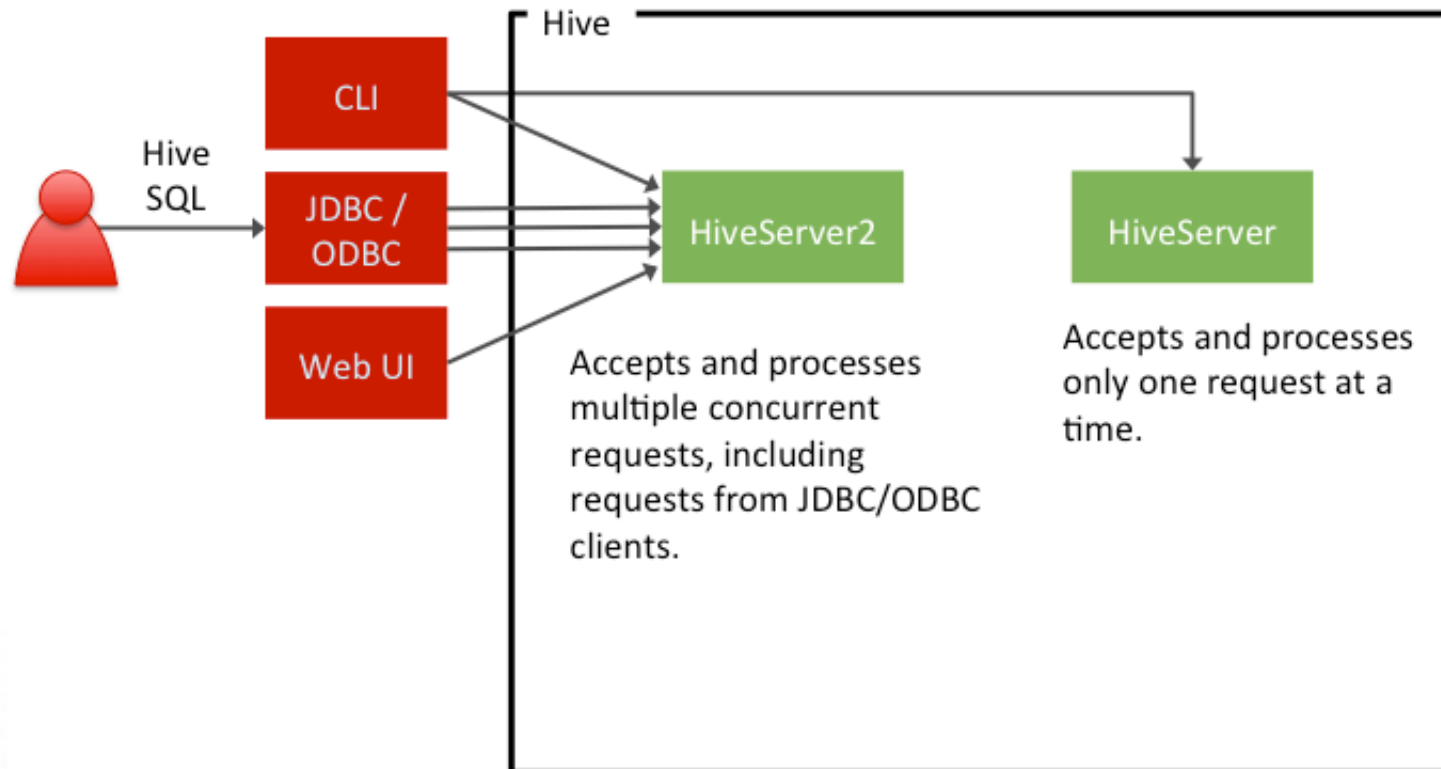
After



Vectorization + ORC files = a huge breakthrough in Hive query performance



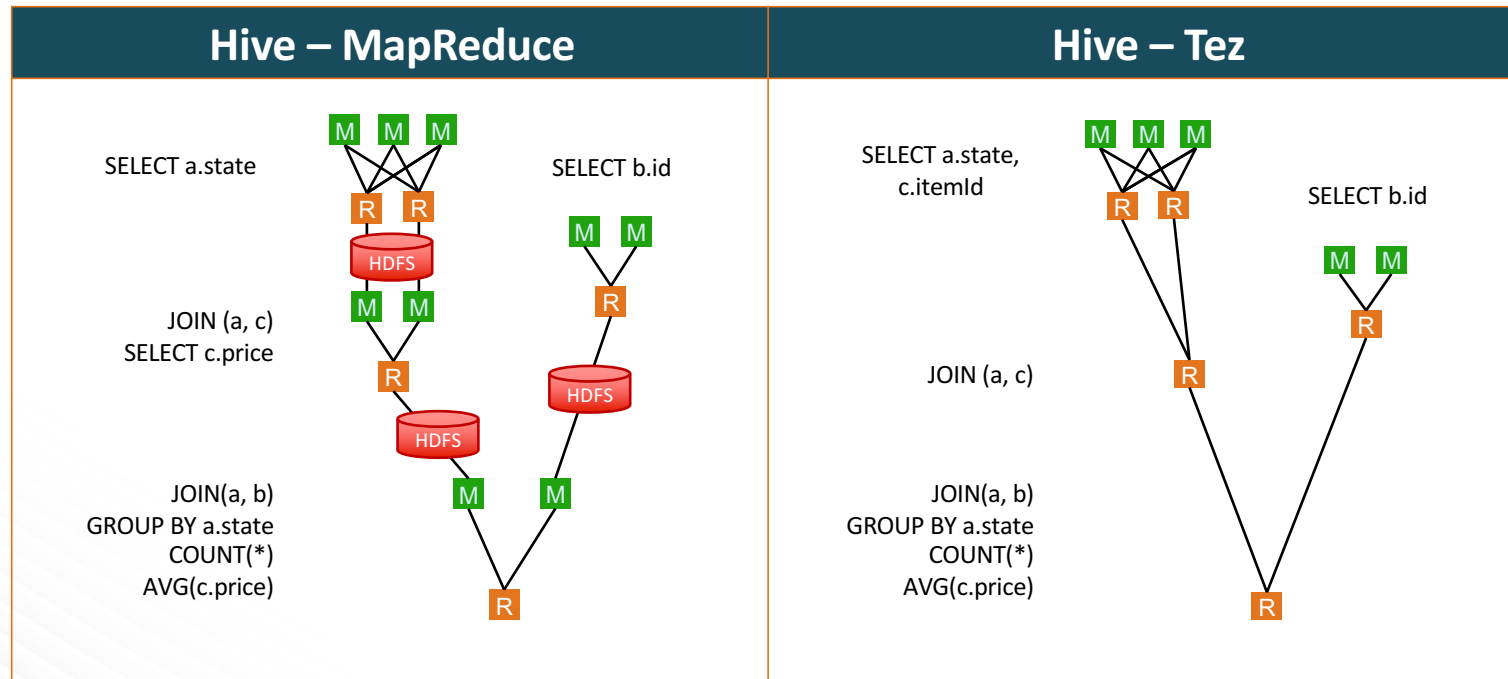
## Using HiveServer2



# Understanding Hive on Tez

```
SELECT a.state, COUNT(*), AVG(c.price)
FROM a
JOIN b ON (a.id = b.id)
JOIN c ON (a.itemId = c.itemId)
GROUP BY a.state
```

Tez avoids unneeded  
writes to HDFS



## Using Tez for Hive Queries

Set the following property in either **hive-site.xml** or in your script:

```
set hive.execution.engine=tez;
```



## Hive Optimization Tips

- ◆ Divide data amongst different files that can be pruned out by using partitions, buckets, and skews
- ◆ Use the ORC file format
- ◆ Sort and Bucket on common join keys
- ◆ Use map (broadcast) joins whenever possible
- ◆ Increase the replication factor for hot data (which reduces latency)
- ◆ Take advantage of Tez



## Hive Query Tunings

- ◆ `mapreduce.input.fileinputformat.split.maxsize`
- ◆ `mapreduce.input.fileinputformat.split.minsize`
- ◆ `mapreduce.tasks.io.sort.mb`

In addition, set the important join and bucket properties to **true** in **hive-site.xml** or by using the **set** command.



## Lesson Review

1. What is the benefit of performing two **insert** queries in the same Hive command?
2. **True or False:** Hive views are materialized when they are defined.
3. Suppose an **employees** table has 200 rows and its **department** column has 15 distinct values. How many rows would be in the result set of the following query?  

```
from employees  
select fname,lname,MAX(salary)  
over (partition by department);
```
4. Explain what the following query is computing:  

```
from employees  
select fname,lname,AVG(salary)  
over (partition by department order by salary  
rows between 5 preceding and current row);
```
5. Which Hive file format provides the best performance?
6. What does DAG stand for?

